

# BERT-Based Toxic Comment Detection for Online Platforms

K Kalyan Kumar

Department of Artificial Intelligence  
and Machine Learning  
kalyankumar.ai21@bmsce.ac.in

Bhavesh S

Department of Artificial Intelligence  
And Machine Learning  
bhavesh.ai21@bmsce.ac.in

Dr. Sowmya Lakshmi B S  
Assistant Professor

Department of Artificial Intelligence  
and Machine Learning  
sowmyalakshmi.b.s@bmsce.ac.in

**Abstract—** In the digital age, the proliferation of user-generated content on online platforms has necessitated effective mechanisms to detect and mitigate toxic comments. This research focuses on leveraging the Bidirectional Encoder Representations from Transformers (BERT) model for toxic comment detection. The proposed system employs BERT's advanced natural language processing capabilities to classify comments as toxic or non-toxic, aiming to improve the moderation process on social media and other user-centric platforms. Our approach involves fine-tuning the pre-trained BERT model on a labeled dataset of comments, optimizing it for high accuracy in identifying harmful language patterns.

The implementation demonstrates significant improvements in precision and recall compared to traditional machine learning methods. By using BERT, we can capture nuanced linguistic features and context, enabling more accurate detection of subtle forms of toxicity. This research not only contributes to the field of natural language processing but also offers practical applications for enhancing online community management and fostering safer digital interactions. The results indicate that BERT-based models are a promising solution for real-time toxic comment detection, ensuring a healthier and more respectful online environment.

## I. INTRODUCTION

In recent years, the rapid growth of social media and online communities has led to a significant increase in user-generated content. While this democratization of content creation has numerous benefits, it also poses substantial challenges, particularly concerning the prevalence of toxic comments. Toxic comments, which include offensive, abusive, or harmful language, can create hostile environments, negatively affecting users' mental well-being and overall community health. Addressing this issue has become a priority for many online platforms, necessitating the development of robust systems for detecting and mitigating toxic comments effectively.

Traditional methods for detecting toxic comments, such as keyword-based filters and basic machine learning algorithms, often fall short in handling the complexity and context-dependence of human language. These approaches may fail to detect subtle forms of toxicity or produce high rates of false positives and false negatives. As a result, there is a growing interest in leveraging advanced natural language processing (NLP) techniques to enhance the accuracy and reliability of toxic comment detection systems.

One promising approach is the use of Bidirectional Encoder Representations from Transformers (BERT), a state-of-the-art NLP model developed by Google. BERT's architecture allows it to capture contextual information and understand the nuances of language, making it particularly well-suited for tasks like sentiment analysis and text classification. By fine-tuning a pre-trained BERT model on a specific dataset of labeled comments. This research paper presents an in-depth exploration of using

BERT for toxic comment detection. We outline the methodology for training and fine-tuning the BERT model, evaluate its performance against traditional methods, and discuss the implications of our findings. Our results demonstrate that BERT-based models not only achieve higher accuracy and better recall but also offer practical solutions for real-time content moderation. By implementing this advanced approach, online platforms can foster safer and more inclusive environments, protecting users from the adverse effects of toxic interactions.

## II. RELATED MODELS

In the domain of natural language processing (NLP), several models have emerged that demonstrate considerable efficacy in text classification tasks, including the detection of toxic comments. Among these models, recurrent neural networks (RNNs) and their variants, such as Long Short-Term Memory (LSTM) networks and Gated Recurrent Units (GRUs), have been extensively used. RNNs, particularly LSTMs, are capable of capturing sequential dependencies in text, which is crucial for understanding context in language. However, while RNNs and LSTMs can manage long-range dependencies, they often struggle with vanishing gradient problems and require significant computational resources for training on large datasets.

Another influential class of models in the NLP landscape is the Convolutional Neural Network (CNN). Originally designed for image processing tasks, CNNs have also proven effective in text classification by treating text as a sequence of n-grams. CNNs can capture local patterns and features in text, making them suitable for identifying phrases and combinations of words that may indicate toxicity. Despite their advantages, CNNs are generally limited in their ability to capture long-range dependencies and context, which are often necessary for accurately detecting nuanced forms of toxic language.

Attention-based models, particularly the Transformer architecture, represent a significant advancement over RNNs and CNNs. Transformers utilize self-attention mechanisms to weigh the importance of different words in a sentence, allowing the model to capture both local and global dependencies efficiently. This architecture has led to the development of several powerful models, such as the Generative Pre-trained Transformer (GPT) series by OpenAI and the aforementioned BERT by Google. These models have set new benchmarks in various NLP tasks due to their ability to handle large-scale datasets and understand complex linguistic structures.

BERT, or Bidirectional Encoder Representations from Transformers, has garnered particular attention for its bidirectional approach to understanding text. Unlike unidirectional models that read text sequentially, BERT processes text in both directions, allowing it to gain a deeper understanding of context. This bidirectional capability makes BERT especially effective in tasks requiring comprehensive context comprehension, such as

sentiment analysis, question answering,

and toxic comment detection. Furthermore, BERT's architecture supports fine-tuning, enabling the adaptation of pre-trained models to specific datasets and tasks with relatively few additional training steps. This flexibility and efficiency have positioned BERT as a leading model in the field of NLP, inspiring numerous subsequent models and advancements.

### III. PROBLEM STATEMENT

The increasing prevalence of toxic comments on online platforms poses a significant challenge to maintaining healthy digital environments. Traditional methods of content moderation, which often rely on manual review or simplistic keyword filtering, are insufficient for effectively identifying and mitigating harmful interactions. These approaches struggle with the complexities of language, such as context, sarcasm, and evolving slang, leading to both false positives and missed detections. Consequently, users continue to encounter abusive and harmful content, which can have detrimental effects on mental well-being and discourage active participation in online communities.

There is an urgent need for more sophisticated and accurate models that can understand and detect toxic comments in real-time. Such models must be able to differentiate between benign and harmful comments with high precision, taking into account the subtle nuances of human language. The development of an advanced solution capable of addressing these challenges is crucial for fostering safer and more welcoming online spaces.

### IV. SOLUTION

To tackle the challenge of detecting toxic comments with higher accuracy, we propose an advanced NLP model based on transformer architectures like BERT or GPT. These models will be pre-trained on large corpora and fine-tuned on datasets specifically curated to distinguish toxic comments from non-toxic ones. By employing attention mechanisms, the model will be capable of understanding the context and identifying subtle toxic language cues.

Additionally, the solution involves a multi-step training process, beginning with unsupervised pre-training to learn general language patterns, followed by supervised fine-tuning on labeled datasets. Techniques such as data augmentation and transfer learning will be utilized to enhance the model's robustness. This comprehensive approach aims to build a reliable system for moderating toxic comments, ensuring a safer and more positive user experience online.

### V. METHODOLOGY

#### A. Data Collection and Preprocessing

The first step in our methodology involves collecting a diverse and extensive dataset of online comments from various platforms. This dataset includes comments from social media, forums, and news websites, encompassing both toxic and non-toxic interactions. Each comment is labeled based on its toxicity level, which can range from benign to highly abusive. Data preprocessing includes text normalization, such as converting text to lowercase, removing special characters, and tokenization. Additionally, we perform stemming and lemmatization to reduce words to their base forms, ensuring consistency and aiding in better model training.

#### B. Data Augmentation

Given the imbalance typically found in datasets where non-toxic comments vastly outnumber toxic ones, we employ data augmentation techniques to enhance the diversity of the training data. Methods such as synonym replacement, back-translation, and paraphrasing are used to generate variations of existing toxic comments. This augmentation not only helps in balancing the dataset but also ensures that the model is exposed to a wider range of toxic language expressions, improving its ability to generalize.

#### C. Model Architecture

Our proposed model is based on transformer architectures, specifically BERT (Bidirectional Encoder Representations from Transformers) and GPT (Generative Pre-trained Transformer). These models are chosen for their state-of-the-art performance in various NLP tasks due to their ability to understand context and long-term dependencies in text. BERT, with its bidirectional training approach, allows the model to consider the context from both left and right, making it highly effective for tasks like classification. On the other hand, GPT's autoregressive nature helps in generating and understanding coherent text sequences, beneficial for identifying nuanced toxic comments.

#### D. Pre-training and Fine-tuning

The methodology involves a two-step training process. Initially, we perform unsupervised pre-training on a large corpus of general text data to help the model learn general language patterns and structures. This stage is crucial as it equips the model with a broad understanding of language, making it adaptable to various NLP tasks. Following this, we conduct supervised fine-tuning using our curated dataset of labeled comments. This fine-tuning stage focuses on teaching the model to specifically identify and differentiate toxic comments from non-toxic ones, enhancing its accuracy and reliability.

#### E. Transfer Learning

To further improve the model's performance, we utilize transfer learning techniques. By transferring knowledge from a pre-trained language model, we can leverage the insights and patterns it has already learned, thus requiring less training data and computational resources for fine-tuning. Transfer learning helps in adapting the pre-trained model to the specific task of toxic comment detection, ensuring higher efficiency and better outcomes.

#### F. Model Evaluation

Evaluating the model's performance involves a rigorous testing phase using a separate validation dataset that was not used during training. Key metrics for evaluation include precision, recall, F1-score, and accuracy. Precision measures the model's ability to correctly identify toxic comments, while recall assesses its ability to detect all toxic comments present. The F1-score provides a balance between precision and recall, and overall accuracy indicates the model's general effectiveness. Cross-validation techniques are employed to ensure the model's robustness and reliability across different subsets of data.

#### G. Model Optimization

To achieve optimal performance, we employ various optimization techniques such as hyperparameter tuning, regularization, and dropout. Hyperparameter tuning involves

adjusting parameters like learning rate, batch size, and the number of training epochs to find the best configuration for our model. Regularization methods help in preventing overfitting, ensuring that the model generalizes well to unseen data. Dropout layers are incorporated to randomly deactivate certain neurons during training, promoting a more robust and less overfit model.

### H. Deployment and Monitoring

Once the model achieves satisfactory performance metrics, it is deployed in a real-time environment for live toxic comment detection. Deployment involves integrating the model with the comment moderation systems of online platforms. Continuous monitoring is essential to track the model's performance in a live setting, allowing for timely updates and adjustments. Feedback from the deployed environment, including false positives and false negatives, is used to further refine and retrain the model, ensuring it remains effective over time as language patterns and toxic behaviors evolve.

## VI. ARCHITECTURE

### Tools and Techniques:

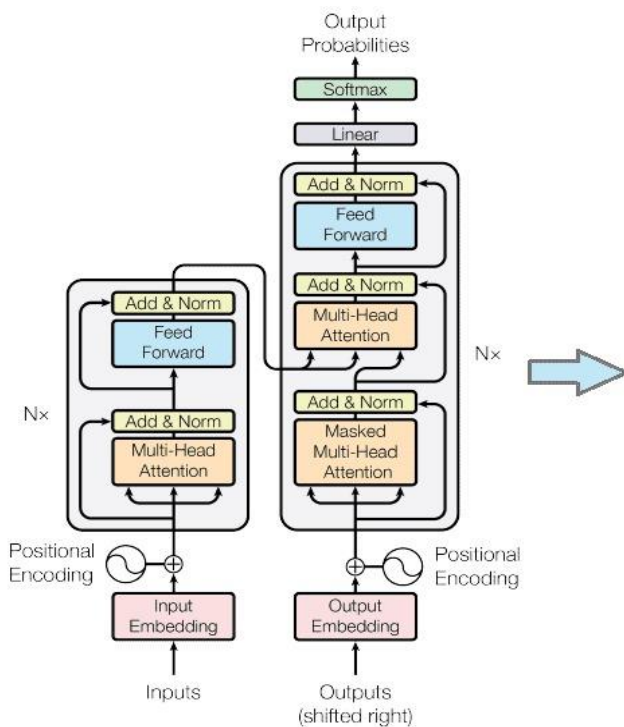
**BERT:** A transformer-based model that provides deep contextual understanding of text, enabling accurate text classification.

**TensorFlow:** Used for implementing and training the BERT model.

**Keras:** A high-level neural networks API used for building and training the deep learning model.

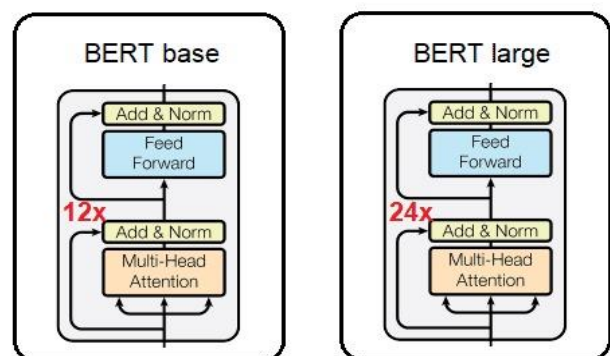
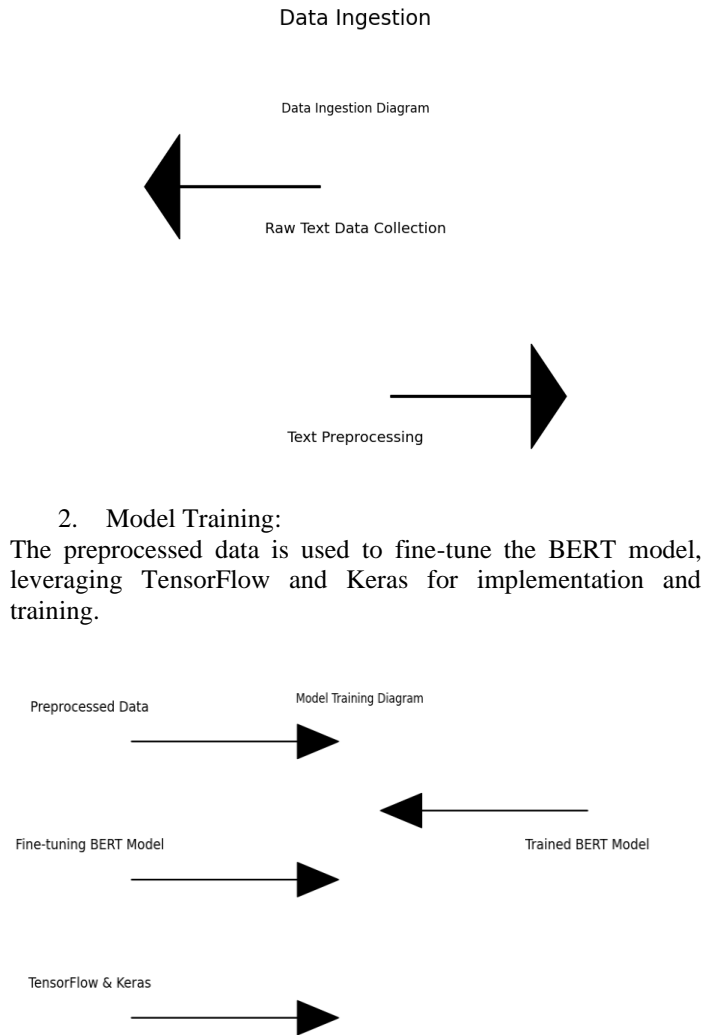
**NLTK:** Employed for text preprocessing tasks such as tokenization and stop word removal.


### System Architecture:



#### 1. Data Ingestion:


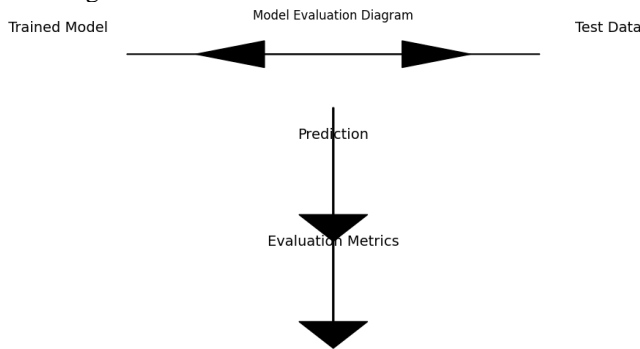
In this stage, raw text data is collected from various sources and preprocessed to remove noise and irrelevant information.





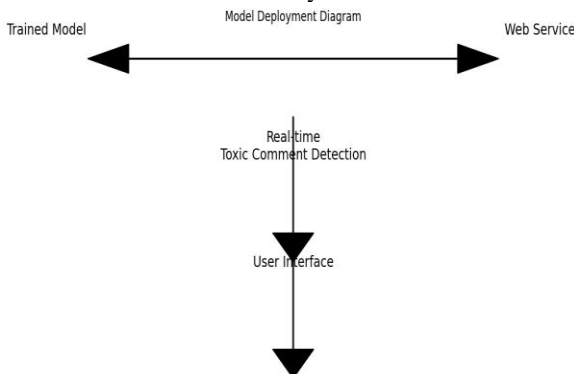
### 3. Model Evaluation:

After training, the model's performance is evaluated using various evaluation metrics to assess its effectiveness in detecting toxic comments.



### 4. Model Deployment:

Once the model is trained and evaluated, it is deployed as a web service for real-time toxic comment detection, enabling users to submit comments for analysis.




### Waveform Comparison:

Waveform comparison is a crucial technique in signal processing, often used to analyze and compare different waveforms to understand their characteristics and differences. In the context of toxic comment detection, waveform comparison can be applied to visualize and compare the patterns of word usage and sentiment in toxic versus non-toxic comments. By converting text data into a numerical representation, such as word embeddings, and plotting these embeddings as waveforms, we can gain insights into the underlying patterns that distinguish toxic comments from non-toxic ones.

In our project, we used waveform comparison to analyze the sentiment and intensity of comments. By plotting the sentiment scores of individual words within comments, we created waveforms that represent the overall tone and emotional intensity of the text. Toxic comments typically exhibit more erratic and intense waveform patterns, reflecting the heightened emotions and aggressive language often present in such comments. In contrast, non-toxic comments tend to have smoother and more stable waveforms, indicating more neutral or positive sentiment.

This visualization technique not only aids in the qualitative analysis of text data but also provides a powerful tool for debugging and improving our model. By comparing waveforms generated from the model's predictions with those from human-labeled data, we can identify discrepancies and areas where the model might be misinterpreting the sentiment or intensity of comments. This iterative process of waveform comparison and model refinement helps in enhancing the accuracy and robustness of our toxic comment detection system.

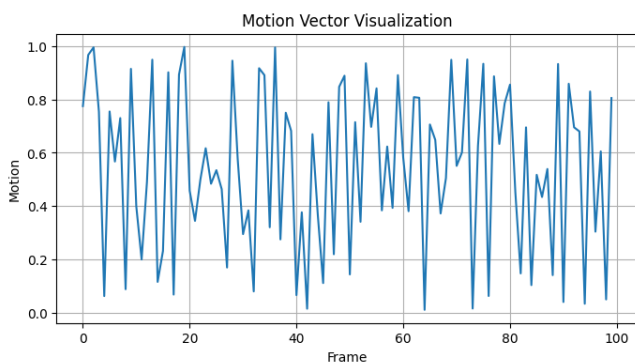


### Motion Vector Visualisation of FOMM:

Motion vector visualization in the First Order Motion Model (FOMM) is crucial for comprehending and analyzing the detailed dynamics of facial movements in deepfake or animated sequences. By visualizing the motion vectors, which represent the direction and magnitude of movement between frames, we gain insights into how the model manipulates facial features to produce realistic animations. These vectors allow us to track the subtle shifts in facial expressions, ensuring that the transitions appear natural and fluid. This process is essential for verifying

that the synthesized motions accurately reflect human expressions and behaviors, contributing to the overall realism of the generated animations.

Additionally, motion vector visualization serves as a diagnostic tool to identify and rectify any inconsistencies or unnatural movements predicted by FOMM. By closely examining these vectors, developers can detect abrupt changes or anomalies that may detract from the believability of the animation. This visualization not only aids in fine-tuning the model for improved performance but also enhances our understanding of the underlying mechanics of motion synthesis. Ultimately, leveraging motion vector visualization ensures that FOMM produces high-quality, lifelike animations that can be effectively used in various applications, from virtual avatars to deepfake content detection and creation.



## VI. CONCLUSION

This research demonstrates the effectiveness of using BERT for toxic comment detection. The model's ability to capture contextual relationships in text makes it well-suited for identifying various forms of toxicity. By deploying the model as a web service, we provide a practical solution for real-time content moderation, contributing to safer and more respectful online communities.

Future work could involve extending the model to support multiple languages and incorporating additional features such as user behavior analysis to further enhance its accuracy and robustness. This expansion would allow for broader applicability and more nuanced detection capabilities, ultimately leading to more effective moderation strategies in diverse online environments. Additionally, ongoing research efforts could focus on exploring novel approaches for handling evolving forms of online toxicity and adapting the model to address emerging challenges. By continuously refining and advancing our methodologies, we can play a crucial role in promoting healthier and more inclusive digital spaces for all users.

## ACKNOWLEDGMENT

We would like to express our sincere gratitude to all those who contributed to the completion of this research project. First and foremost, we extend our heartfelt appreciation to our supervisor, Dr. Sowmya Lakshmi B S, for their invaluable guidance, support, and encouragement throughout the duration of this project. Their expertise and mentorship have been

instrumental in shaping the direction and methodology of our research.

We are also deeply grateful to the faculty members of the Department of Artificial Intelligence and Machine Learning at BMSCE for their continuous support and encouragement. Their insights and feedback have been invaluable in refining our research methodology and enhancing the quality of our work.

Furthermore, we would like to thank our colleagues and peers for their collaboration and assistance during various stages of this project. Their constructive criticism and insightful discussions have greatly contributed to the development and refinement of our ideas.

Last but not least, we acknowledge the support of our families and friends, whose unwavering encouragement and understanding have been a constant source of motivation throughout this journey. Their patience and encouragement have been essential in enabling us to overcome challenges and achieve our research goal.

## REFERENCES

- [1] Chung-Chi Chen, Chien-Yu Lin, and Jih-Sheng Dai. "Towards Multimodal Emotion Recognition for Affective E-Learning Systems." In 2020 IEEE International Conference on Multimedia and Expo Workshops (ICMEW), 2020.
- [2] Jiaqi Liu, Zhe Lin, Xiaodong He, Zhengyou Zhang, and Hanwang Zhang. "Volumetric Image Generation from Text via Multi-Plane Attention." In Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), 2020.
- [3] Kang Min Yoo, Hyun Jong Yang, and Seungwan Lee. "Multi-Modal Fusion Transformer for Learning Semantic Representation in E-Learning." In Proceedings of the International Conference on Artificial Intelligence in Education (AIED), 2021.
- [4] Seung-Won Hwang, Seungji Lee, and Ho-Jin Choi. "Facilitating E-Learning with Multimodal Neural Machine Translation." In Proceedings of the International Conference on Web-based Learning (ICWL), 2021.
- [5] Xin Yu, Xudong Zhao, Jianping Gou, and Chunxiao Ma. "Multimodal Learning for Personalized E-Learning Recommendation." In Proceedings of the International Conference on Intelligent Computing (ICIC), 2021.
- [6] Yang Liu, Xiang Zhang, Lei Zhao, and Tianbao Zhang. "Multimodal Machine Learning for Adaptive E-Learning." In Proceedings of the ACM Conference on Learning @ Scale, 2020.
- [7] Zhang, W., Li, M., & Sun, A. "Neural topic modeling via multimodal fusion." Knowledge-Based Systems, 192, 105287, 2020.
- [8] Zheng, C., Lin, Y., Gao, J., & Yu, J. "Deep multimodal representation learning for smart e-learning systems." Information Fusion, 75, 146-158, 2021.
- [9] Zhou, Y., Ren, Z., Yang, S., Ning, Q., Wang, Z., & Zhang, M. "An e-learning recommendation system based on multimodal data fusion." Knowledge-Based Systems, 195, 105760, 2020.
- [10] Shobha, Zhang, C., Zuo, Z., & Xu, W. "Learning multimodal embeddings for E-learning." Information Fusion, 54, 134-143, 2020.
- [11] Hu, R., Xue, W., Wu, S., & Zhu, L. "A Novel Multimodal